

TECHNICAL SEO

Adhiroha — Action Plan

adhiroha.com · Next.js 15 App Router · 11 languages, 213 pages

66

OF 100

Audit based on source plus the 213 pre-rendered HTML files in the current build — the same HTML Googlebot receives. Two items carry real risk; the rest is performance work and polish.

Prepared 30 July 2026

Method Static analysis of source + build output. No live HTTP requests were made.

Not verified Live SSL, redirects, response headers as served, CrUX field data

Where the site stands

CATEGORY	STATUS	SCORE	HEADLINE
Crawlability	pass	88	robots + sitemap correct; 211 routes covered
Indexability	warn	72	hreflang exemplary; lang and Twitter tags wrong
Security	warn	55	no security headers configured
URL structure	pass	82	clean and consistent; legacy footer links leak
Mobile	pass	78	responsive, CLS well protected
Core Web Vitals	fail	35	render-blocking fonts, Leaflet, Hotjar
Structured data	warn	60	rich graph, but self-serving ratings
JS rendering	pass	96	everything server-rendered
IndexNow	fail	0	not implemented

One correction to the earlier read. The `<html lang="en">` issue was first flagged as high priority. Google ignores document lang attributes for language detection — it reads visible content. So this is a genuine defect for screen readers and for Bing, but not a Google ranking issue. It has been moved to Phase 2.

Phase 0 — This week

≈2–3 hours · ~85% of the value

0.1 · Remove the self-serving AggregateRating

Critical

20 min

lib/seo.js

Every page emits `aggregateRating: 5.0, reviewCount: 300` in two places — on `EducationalOrganization / LocalBusiness` and on every `Course`. There are no `Review` nodes and no third-party source, and the value is identical on all 213 pages.

Google's rich-result policy disallows self-serving ratings for local businesses and organisations: ratings must come from users, not be compiled by the site. A flat 5.0 with a round 300 and no review data is a well-known trigger for a structured-data manual action.

Do: delete the `aggregateRating` block from `organizationSchema()` and `courseSchema()`. Star ratings in search results will be lost — but they were always at risk, and a manual action would cost far more.

To earn them back properly: embed a Google Business Profile review widget that emits its own schema, or add real `Review` nodes with genuine `author`, `datePublished` and `reviewBody`.

0.2 · Take Leaflet off the critical path

Largest CWV win

1 hour

app/layout.jsx

The root layout loads Leaflet on **all 213 pages**: a render-blocking stylesheet in `<head>` and a script with `strategy="beforeInteractive"`, which delays interactivity everywhere.

A map is only present on **111 pages** (those with a `#gmap` container), and it always sits just above the footer — far below the fold. So 102 pages never need it at all, and the other 111 do not need it immediately.

Do: remove both from the root layout and inject them on demand when the map approaches the viewport.

```
// components/LeafletOnDemand.jsx
"use client";
import { useEffect } from "react";

// Leaflet is only needed by the map that sits just above the footer, so it must
// never block first paint. We inject the CSS+JS the first time that map comes
// near the viewport, then let the page's own init script run.
export default function LeafletOnDemand() {
  useEffect(() => {
    const el = document.getElementById("gmap");
    if (!el || window.L) return;

    const load = () => {
      const css = document.createElement("link");
      css.rel = "stylesheet";
      css.href = "https://unpkg.com/leaflet@1.9.4/dist/leaflet.css";
      css.crossOrigin = "";
      document.head.appendChild(css);

      const js = document.createElement("script");
      js.src = "https://unpkg.com/leaflet@1.9.4/dist/leaflet.js";
      js.crossOrigin = "";
      js.onload = () => window.dispatchEvent(new Event("leaflet:ready"));
      document.head.appendChild(js);
    };

    const io = new IntersectionObserver([e] => {
      if (e.isIntersecting) { io.disconnect(); load(); }
    }, { rootMargin: "400px" });
    io.observe(el);
    return () => io.disconnect();
  }, []);
  return null;
}
```

Each page's `scripts.js` must then wait for `leaflet:ready` before initialising its map. Because those script files are duplicated per page, apply that change with a single script across all of them — the same approach used for the translations.

Bonus: self-host Leaflet in `public/js/` to drop the third-party DNS lookup and TLS handshake as well.

0.3 · Defer Hotjar

5 min

app/layout.jsx

Hotjar is inlined in `<head>`, so it executes before first paint and adds to total blocking time on every page. Move it to the end of `<body>` and let it load when the page is idle — heatmaps and recordings are unaffected.

```
<Script id="hotjar" strategy="lazyOnload"
  dangerouslySetInnerHTML={{ __html: HOTJAR_SCRIPT }} />
```

0.4 · Fix the legacy non-www footer links

20 min

209 pages

Three legacy URLs sit in the server-rendered footer of 209 pages. Each is non-www, so it costs a 301 to www and then a 308 for the trailing slash — a two-hop chain, site-wide.

`components/FooterLinkFix.jsx` rewrites them, but only inside `useEffect` — after hydration. The HTML Googlebot first receives still carries the wrong URLs.

Do: correct them in the source content with an assert-based replacement script, then delete `FooterLinkFix` entirely.

```
const FIX = {
  "https://adhiroha.com/weather.php": "/weather/",
  "https://www.adhiroha.com/weather.php": "/weather/",
  "https://adhiroha.com/volunteer-opportunity-in-rishikesh": "/volunteer-opportunity-
in-rishikesh/",
  "https://www.adhiroha.com/volunteer-opportunity-in-rishikesh": "/volunteer-
opportunity-in-rishikesh/",
  "https://adhiroha.com/apply-for-teacher-in-rishikesh": "/apply-for-teacher-in-
rishikesh/",
  "https://www.adhiroha.com/apply-for-teacher-in-rishikesh": "/apply-for-teacher-in-
rishikesh/",
};
```

Phase 1 — Next one to two weeks

performance

1.1 · Pull the fonts out of the CSS

Biggest single CWV item

half a day

The route stylesheet is 692 KB, of which 418 KB is base64 font binary across eight `@font-face` blocks. The browser cannot paint any text until all of it has downloaded.

Do: extract the `.woff2` files to `public/fonts/`, reference them with `url()`, and `<link rel="preload">` the two or three faces used above the fold. Render-blocking CSS drops to roughly 274 KB and the fonts load in parallel.

This is item P0.1 of the existing `MOBILE_PAGESPEED_ACTION_PLAN.md`, and that analysis is sound.

1.2 · Images — lazy-load and compress

2 hours

Only **11 of 56** images carry `loading="lazy"`, so 45 load eagerly. (The existing plan states 46 of 56 — that figure is wrong; 11 was measured directly from the build.)

- Add `loading="lazy" decoding="async"` to everything except the hero — scripted, one pass.
- Compress the outliers: `favicon.png` 1 MB, `background-1.png` 2.3 MB, `om-logo.png` 1.4 MB.

Note `next.config.mjs` sets `images: { unoptimized: true }`, so Next does no resizing or format conversion — every byte shipped is a byte authored.

1.3 · Twitter Card metadata on translated pages

30 min

184 files

184 language `page.jsx` files set `openGraph` but not `twitter`, so the root layout's English Twitter card is inherited. A Danish page shared on X shows an English title.

Do: each file already holds a translated `openGraph.title` and `description` — derive `twitter` from them with a script. Purely mechanical.

Phase 2 — Polish and backlog

2.1 · Security headers

15 min

`next.config.mjs`

No headers are configured at all. Ranking impact is small — only HTTPS is a confirmed signal, and a light one — but this is standard hygiene and shows up in Best Practices audits.

```
async headers() {
  return [{
    source: "/*:path*",
    headers: [
      { key: "Strict-Transport-Security", value: "max-age=63072000; includeSubDomains; preload" },
      { key: "X-Content-Type-Options", value: "nosniff" },
      { key: "Referrer-Policy", value: "strict-origin-when-cross-origin" },
      { key: "X-Frame-Options", value: "SAMEORIGIN" },
    ],
  }];
}
```

Leave CSP for later — Hotjar, Leaflet, Razorpay and the inline scripts all need allowing, and that deserves its own careful pass.

2.2 · `<html lang>` on translated pages

Accessibility + Bing

30 min

All 200+ localised pages serve `<html lang="en">` while their canonical and hreflang say otherwise. Google ignores this attribute, but screen readers will read Danish text in an English voice, and Bing does use it.

In the App Router only the root layout may render `<html>`, and it cannot read the pathname in a static build. Two options:

- **Cheap (recommended):** a small client component that sets `document.documentElement.lang` — one line per language layout. Screen readers read the live DOM and Bing renders JS, so both are covered.
- **Thorough (optional):** drop `app/layout.jsx` and give each language its own root layout — 11 layouts sharing a common head component. Fully static, but a sizeable refactor for a signal Google does not read.

2.3 · AI crawler policy — a decision, then five minutes

Business decision

`app/robots.js` allows everything, so GPTBot, ClaudeBot, PerplexityBot, Google-Extended and CCBot all have access. For a yoga school the recommendation is to **keep it that way** — citations in ChatGPT and Perplexity drive brand awareness and referral traffic, and the content is genuinely useful.

If training use should be limited, blocking `Google-Extended` is safe: it stops Gemini training but has no effect on Google Search or AI Overviews, which use `Googlebot`.

2.4 · IndexNow

30 min

Low impact

Add a key file at `public/<key>.txt` and ping changed URLs after each deploy. Since articles are database-driven, new posts would reach Bing, Yandex and Naver quickly. Google does not support IndexNow, so the benefit is limited to those engines.

2.5 · Small items

Backlog

- `/student-admission-panel/` is indexable but absent from the sitemap — pick one. (`/thank-you/` is correctly `noindex, follow`.)
- 38 CSS rules set 11–13px type; worth a readability check on mobile.
- `llms.txt` is absent — an opportunity, not a defect. Lighthouse's new Agentic Browsing category looks for it.

Worth protecting

These are already right, and several are things most multilingual sites get wrong.

- **Hreflang is exemplary** — 11 languages plus `x-default` , reciprocal, absolute URLs, on every page.
- **Everything is server-rendered** — content, canonical, meta robots and JSON-LD all appear in the initial HTML, matching Google's December 2025 JavaScript SEO guidance.
- **Page weight is healthy** — the largest page is 452 KB, far inside Googlebot's 2 MB HTML limit.
- **CLS is well protected** — 331 explicit `width / height` attributes.
- **Fonts are self-hosted** with `font-display: swap` ; no external font requests.
- **URLs are clean** — descriptive, hyphenated, consistent trailing slashes, longest is 88 characters.
- **The back button is safe** — `AnchorScroll` uses `replaceState` rather than `pushState` , so in-page anchors never trap history.

Still to verify live

Documented in `DEPLOY.md` but not confirmed, because no network requests were made. If any of these are broken they outrank everything in Phase 0.

- Is the non-www → www 301 actually live?
- Valid SSL certificate, and no mixed content?
- Do the legacy `*.html` redirects still work?
- Real CrUX field data for LCP, INP and CLS.

Prepared 30 July 2026 · Scores are weighted by search impact, with Core Web Vitals and indexability carrying the most weight and IndexNow the least. All findings were verified against the build output rather than inferred from source alone.